


In [85]:

```
1 #Loading required libraries
2 import numpy as np
3 import pandas as pd
4 import matplotlib.pyplot as plt
5 from sklearn.linear_model import LinearRegression
6 from sklearn.metrics import mean_squared_error
7 from sklearn.impute import SimpleImputer
8 from sklearn.ensemble import HistGradientBoostingRegressor
9 import matplotlib.pyplot as plt
10 import seaborn as sns
11 from sklearn.model_selection import cross_val_score
12 import seaborn as sns
13 import statsmodels.api as sm
14
```

In [86]:

```
1 train_data = pd.read_csv('train.csv')
2 test_data=pd.read_csv('test.csv')
3 data=pd.concat([train_data,test_data])
4 plt.figure(figsize=(18,6))
5 plt.title('Heatmap of missing values')
6 sns.heatmap(data.isnull(),yticklabels=False,cbar=False,cmap='viridis')
7
```

Out[86]: <Axes: title={'center': 'Heatmap of missing values'}>



```
In [87]: 1 # separate numerical and categorical columns
2 numerical_cols = train_data.select_dtypes(include=['int64', 'float64'])
3 categorical_cols = train_data.select_dtypes(include=['object']).columns
4
5 # replace missing values with mean for numerical columns
6 imputer = SimpleImputer(strategy='mean')
7 train_data[numerical_cols] = imputer.fit_transform(train_data[numerical_
8
9 # replace missing values with mode for categorical columns
10 imputer = SimpleImputer(strategy='most_frequent')
11 train_data[categorical_cols] = imputer.fit_transform(train_data[categor:
12
13 # verify that missing values are replaced
14 print(train_data.isnull().sum())
15
16
17
```

```
Id          0
MSSubClass  0
MSZoning    0
LotFrontage 0
LotArea     0
..
MoSold      0
YrSold      0
SaleType    0
SaleCondition 0
SalePrice   0
Length: 81, dtype: int64
```

```
In [88]: 1 # Select the predictor and target variable (for both training and testing)
2 x_train = train_data[['LotArea', 'BedroomAbvGr', 'BsmtFullBath', 'BsmtHalfBath']]
3 y_train = train_data['SalePrice']
4
5 x_test = test_data[['LotArea', 'BedroomAbvGr', 'BsmtFullBath', 'BsmtHalfBath']]
```

```
In [89]: 1 # separate numerical and categorical columns for test data
2 numerical_cols2 = x_test.select_dtypes(include=['int64', 'float64']).columns
3
4 # replace missing values with mean for numerical columns
5 imputer = SimpleImputer(strategy='mean')
6 x_test[numerical_cols2] = imputer.fit_transform(x_test[numerical_cols2])
7
8 # verify that missing values are replaced
9 print(x_test.isnull().sum())
```

```
LotArea      0
BedroomAbvGr 0
BsmtFullBath 0
BsmtHalfBath 0
FullBath      0
HalfBath      0
dtype: int64
```

C:\Users\Lenovo\AppData\Local\Temp\ipykernel_8988\281113641.py:6: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
x_test[numerical_cols2] = imputer.fit_transform(x_test[numerical_cols2])
```

```
In [90]: 1 # Initializing a linear regression model and fit this model to train set
2 model = LinearRegression()
3 print(model)
4 model.fit(x_train, y_train)
5
```

```
LinearRegression()
```

```
Out[90]: ▾ LinearRegression
LinearRegression()
```

```
In [91]: 1 # Making predictions on the test set using the fitted model
2
3 y_pred = model.predict(x_test)
4 print(y_pred)
5
```


```
[112606.119798  145550.48551853 224695.03092123 ... 146582.41859758
120117.45460539 219383.04414268]
```

```
In [92]: 1 #instead of simple linear regression we are using HistGradientBoosting v
2 #missing values are assigned to the left or right child accordingly1. Th
```

```
In [81]: 1 train_data = pd.read_csv('train.csv')
2 test_data=pd.read_csv('test.csv')
3
4 x_train = train_data[['LotArea', 'BedroomAbvGr', 'BsmtFullBath','BsmtHa
5 y_train = train_data['SalePrice']
6
7 x_test=test_data[['LotArea', 'BedroomAbvGr', 'BsmtFullBath','BsmtHalfBa
8 # Initializing a linear regression model and fit this model to train set
9 histgrad_model = HistGradientBoostingRegressor().fit(x_train, y_train)
10
11 # Making predictions on the test set using the fitted model
12 y_pred = histgrad_model.predict(x_test)
13 print("predicted saleprices;",y_pred)
14
```

```
predicted saleprices; [129151.88642225 178516.40056678 253716.44635059 ...
199879.47712159
137533.8702215 208269.17549794]
```

```
In [82]: 1 plt.figure()
2 plt.title('Comparison of Predicted Sale Price and Original Price')
3 plt.scatter(y_train, model.predict(x_train), label='Simple Linear Regres
4 plt.legend()
5 plt.show()
6
```



```
In [93]: 1 plt.figure()
2 plt.title('Comparison of Predicted Sale Price and Original Price')
3 plt.scatter(y_train, histgrad_model.predict(x_train), label='HistGradient')
4 plt.legend()
5 plt.show()
6
```



```
In [96]: 1 #preparing the submission data
2 y_pred = model.predict(x_test)
3 sub_data=pd.DataFrame()
4 sub_data['Id']=test_data['Id']
5 sub_data['SalePrice']=y_pred
6 sub_data
7 sub_data.to_excel('submission.xlsx', index=False)
8
```

In [97]:

```
1  ###LINEAR REGRESSION USING ORDINARY LEAST SQUARE METHOD
2
3  data = pd.read_csv('train.csv')
4
5  # Select the relevant features: square footage, number of bedrooms, and
6  X1 = data[['LotArea', 'BedroomAbvGr', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath']]
7  y1 = data['SalePrice'] # Target variable (house price)
8
9  # Add a constant term for the intercept in the regression model
10 X1 = sm.add_constant(X1)
11
12 # Fit the linear regression model
13 model = sm.OLS(y1, X1).fit()
14 model
15 # Print the summary of the model
16 print(model.summary())
17
```

OLS Regression Results

```

=====
===
Dep. Variable:          SalePrice    R-squared:                0.
459
Model:                  OLS          Adj. R-squared:            0.
457
Method:                 Least Squares    F-statistic:              20
5.9
Date:                  Tue, 04 Jun 2024    Prob (F-statistic):       4.28e-
190
Time:                  03:01:15          Log-Likelihood:           -180
95.
No. Observations:      1460            AIC:                     3.620e
+04
Df Residuals:          1453            BIC:                     3.624e
+04
Df Model:              6
Covariance Type:       nonrobust
=====

```

```

=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
const          3.178e+04    6555.902      4.847      0.000      1.89e+04      4.4
6e+04
LotArea         1.2639        0.158      7.981      0.000          0.953
1.575
BedroomAbvGr -6777.5551     2087.330     -3.247      0.001     -1.09e+04     -268
3.053
BsmtFullBath   3.694e+04     3074.859     12.014      0.000      3.09e+04      4.
3e+04
BsmtHalfBath   1.578e+04     6532.302      2.416      0.016      2967.754      2.8
6e+04
FullBath        7.97e+04     3014.745     26.436      0.000      7.38e+04      8.5
6e+04
HalfBath        3.638e+04     3134.939     11.604      0.000      3.02e+04      4.2
5e+04
=====
=====

```

```

===
Omnibus:          539.540    Durbin-Watson:          1.
975
Prob(Omnibus):    0.000    Jarque-Bera (JB):        3595.
355
Skew:             1.564    Prob(JB):
0.00
Kurtosis:         10.023    Cond. No.                6.70e
+04
=====
=====

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 6.7e+04. This might indicate that there are strong multicollinearity or other numerical problems.

In []:

1